


Optimize Manufacturing Operations with a Predictive Maintenance Model

1. Data Exploration & Validation Strategy

- Plot all sensor readings (temperature, vibration, pressure, current, etc.) over time and overlay known failure events.
 - This helps visually identify patterns, anomalies, and pre-failure trends.
- Check for missing timestamps, duplicates, sudden jumps, or sensor drift.
- Study the temporal behavior of machines (daily patterns, shift-wise variations, seasonal effects).
- Use a time-aware validation strategy — traditional k-fold must not be used.
 - Apply TimeSeriesSplit, rolling-forward validation, or expanding-window validation to simulate real-world model deployment.
- Ensure that validation data always occurs later in time than training data to prevent data leakage.



2. Feature Engineering

- Create features that capture trends, volatility, and time-based patterns in signals.
- Strong features commonly include:
 - Rolling statistics:
 - Rolling mean, rolling std deviation
 - Rolling min/max
 - Rolling median
 - Useful windows: 1 hr, 3 hr, 6 hr, 12 hr
 - Exponential moving averages (EMA) to smooth noisy sensor readings.
 - Gradient / Rate-of-change:
 - First derivative, second derivative
 - Helps detect increasing vibration/temperature before failure
 - Operational features:
 - Time since last maintenance
 - Machine age
 - Shift number (morning/evening/night)
- Capture consistent lookback windows so the model sees fixed-size feature sets for each timestamp.

```
[49]: df['rolling_std_3'] = df['sensor_reading'].rolling(window=3).std().bfill()
df['ema_5'] = df['sensor_reading'].ewm(span=5, adjust=False).mean()
df['gradient'] = np.gradient(df['sensor_reading'])
df['time_since_maintenance'] = np.random.randint(1, 100, size=len(df))
df['operational_age'] = (df['timestamp'] - df['timestamp'].min()).dt.total_seconds() / 3600

print(df.head(10))
```

	timestamp	sensor_reading	failure_event	rolling_std_3	ema_5	gradient	time_since_maintenance	operational_age
0	2023-01-01 00:00:00	58.820262	0	3.422770	58.820262	-6.819476	91	0.0
1	2023-01-01 01:00:00	52.000786	0	3.422770	56.547103	-1.963286	73	1.0
2	2023-01-01 02:00:00	54.893690	0	3.422770	55.995965	4.601840	79	2.0
3	2023-01-01 03:00:00	61.204466	1	4.706423	57.732132	2.222050	36	3.0
4	2023-01-01 04:00:00	59.337790	0	3.241923	58.267352	-8.045428	77	4.0
5	2023-01-01 05:00:00	45.113611	1	8.800827	53.882771	-2.293674	82	5.0
6	2023-01-01 06:00:00	54.750442	0	7.259931	54.171995	2.064802	88	6.0
7	2023-01-01 07:00:00	49.243214	0	4.834799	52.529068	-2.633268	82	7.0
8	2023-01-01 08:00:00	49.483906	0	3.112445	51.514014	1.404889	97	8.0
9	2023-01-01 09:00:00	52.052993	0	1.557401	51.693673	0.618156	50	9.0



3. Modeling for Imbalanced Data

- Failure data is naturally imbalanced; only a small percentage of timestamps correspond to failures.
- Apply class imbalance handling:
 - Use `scale_pos_weight` in XGBoost/LightGBM
 - Or `class_weight='balanced'` in RandomForest or Logistic Regression
- Optionally test oversampling methods like SMOTE — only inside training folds, never across time boundaries.
- Primary evaluation metrics should be:
 - Precision-Recall Curve
 - F1-score
 - PR-AUC (Precision–Recall AUC)
- Avoid accuracy as it gives misleading results for imbalanced problems.
- Tune decision thresholds for your operational needs (high recall for safety, high precision for reducing false alarms).

Fold 1 Classification Report:				
	precision	recall	f1-score	support
0	0.85	0.23	0.36	154
1	0.05	0.50	0.09	12
accuracy			0.25	166
macro avg	0.45	0.36	0.22	166
weighted avg	0.80	0.25	0.34	166



Fold 1 Average Precision (AP): 0.0808

Fold 1 F1-score: 0.0876





Fold 2 Classification Report:				
	precision	recall	f1-score	support
0	0.93	0.27	0.42	158
1	0.04	0.62	0.08	8
accuracy			0.29	166
macro avg	0.49	0.45	0.25	166
weighted avg	0.89	0.29	0.41	166

Fold 2 Average Precision (AP): 0.0405
 Fold 2 F1-score: 0.0781



Fold 3 Classification Report:				
	precision	recall	f1-score	support
0	1.00	0.12	0.22	157
1	0.06	1.00	0.12	9
accuracy			0.17	166
macro avg	0.53	0.56	0.17	166
weighted avg	0.95	0.17	0.21	166

Fold 3 Average Precision (AP): 0.0522
 Fold 3 F1-score: 0.1154





Fold 4 Classification Report:

	precision	recall	f1-score	support
0	0.91	0.13	0.22	159
1	0.03	0.71	0.07	7

	accuracy	macro avg	weighted avg
accuracy	0.47	0.42	0.14
macro avg	0.87	0.15	0.21
weighted avg	0.47	0.42	0.14

Fold 4 Average Precision (AP): 0.0547
Fold 4 F1-score: 0.0662






Fold 5 Classification Report:

	precision	recall	f1-score	support
0	0.91	0.21	0.34	152
1	0.08	0.79	0.15	14

	accuracy	macro avg	weighted avg
accuracy	0.50	0.50	0.25
macro avg	0.84	0.26	0.33
weighted avg	0.50	0.50	0.25

Fold 5 Average Precision (AP): 0.1541
Fold 5 F1-score: 0.1517



```

tscv = TimeSeriesSplit(n_splits=5)

print("TimeSeriesSplit Ready!")

Dataset Loaded Successfully!
   UDI Product ID Type Air temperature [K] Process temperature [K] \
0    1      M14860    M        298.1            308.6
1    2      L47181    L        298.2            308.7
2    3      L47182    L        298.1            308.5
3    4      L47183    L        298.2            308.6
4    5      L47184    L        298.2            308.7

   Rotational speed [rpm] Torque [Nm] Tool wear [min] Machine failure TWF \
0             1551     42.8           0            0       0
1             1408     46.3           3            0       0
2             1498     49.4           5            0       0
3             1433     39.5           7            0       0
4             1408     40.0           9            0       0

   HDF PWF OSF RNF
0   0   0   0   0
1   0   0   0   0
2   0   0   0   0
3   0   0   0   0
4   0   0   0   0
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column          Non-Null Count  Dtype  
---  --  
 0   UDI             10000 non-null   int64  
 1   Product ID      10000 non-null   object  
 2   Type            10000 non-null   object  
 3   Air temperature [K] 10000 non-null   float64 
 4   Process temperature [K] 10000 non-null   float64 
 5   Rotational speed [rpm] 10000 non-null   int64  
 6   Torque [Nm]      10000 non-null   float64 
 7   Tool wear [min]  10000 non-null   int64  
 8   Machine failure 10000 non-null   int64  
 9   TWF             10000 non-null   int64  
 10  HDF             10000 non-null   int64  
 11  PWF             10000 non-null   int64  
 12  OSF             10000 non-null   int64  
 13  RNF             10000 non-null   int64  
dtypes: float64(3), int64(9), object(2)
memory usage: 1.1+ MB
None
Preprocessing Complete! Encoded + Scaled shape: (10000, 10014)
Positive: 339 Negative: 9661 scale_pos_weight: 28.49852507374631
TimeSeriesSplit Ready!

```

4. Dashboard Development

- Design the dashboard with real end-users in mind — typically maintenance engineers.
- Main dashboard should include:
 - Total machine count, failure count, failure rate
 - High-risk machine list
 - Machine selector (dropdown)
- For each selected machine, show:
 - Historical sensor readings with failure markers
 - Predicted risk score (numeric or gauge indicator)
 - SHAP explanation for contributing factors
- Add interactive elements:
 - Zoomable sensor charts
 - Hover tooltips
 - Dynamic risk trend plots
- Ensure the UI is responsive and guides the user quickly to “at-risk” assets.

Predictive Maintenance Dashboard

Total Machines	Failures (Last 100 Days)	Failure Rate (%)
100	196	1.96%

Failure rate calculated over the last 100 days of sensor data

Select Machine from Top 10 High-Risk Machines


Select Machine UDI (Top 10 High-Risk Machines):

4

Or select Machine UDI (Full List):

1


Sensor History for Machine UDI 4




Failure Status: No Failure

Risk Score: 1.00 (High Risk)

Machine Risk Gauge



Important Factors Influencing Risk



Tool wear

0.98

Process temperature

0.03


Rotational speed

0.64


Torque

0.24

Monthly Failure Trend (Machine UDI 4)



SHAP Explanation for Machine UDI 4



[Download Selected Machine Data as CSV](#)

5. Project Structure & Reproducibility

- Organize the project clearly to ensure repeatability:
 - Separate directories for raw data, processed data, notebooks, scripts, and dashboards.
- Maintain a requirements.txt file to fix library versions.
- Use consistent random seeds for reproducibility in all notebooks and model scripts.
- Maintain a clear folder for experiment logs, evaluation metrics, and model artifacts.
- Export the final trained model and ensure the dashboard loads it correctly each time.

```
altair==5.5.0
annotated-doc==0.0.4
anyio==4.12.0
blinker==1.9.0
cachetools==5.5.2
cloudpickle==3.1.2
contourpy==1.3.1
cycler==0.12.1
fastapi==0.123.0
fonttools==4.55.3
gitdb==4.0.12
GitPython==3.1.44
h11==0.16.0
Jinja2==3.1.6
joblib==1.4.2
jsonpatch==1.32
jsonpointer==2.1
jsonschema==4.19.0
kiwisolver==1.4.7
llvmlite==0.46.0
MarkupSafe==3.0.2
matplotlib==3.10.0
narwhals==1.31.0
nbformat==5.8.0
numba==0.63.1
numpy==2.2.1
pandas==2.2.3
plotly==6.5.0
protobuf==5.29.4
PyAudio==0.2.14
pycryptodome==3.23.0
pydeck==0.9.1
Pygments==2.16.1
pyotp==2.9.0
pyparsing==3.2.0
PyQt5==5.15.10
PyQtWebEngine==5.15.6
pywin32==305.1
PyYAML==6.0
scikit-learn==1.6.0
scipy==1.14.1
setuptools==75.1.0
shap==0.50.0
slicer==0.0.8
SpeechRecognition==3.14.
starlette==0.50.0
streamlit==1.43.2
tenacity==9.0.0
threadpoolctl==3.5.0
toml==0.10.2
tornado==6.4.2
traitlets==5.10.1
tzdata==2024.2
ujson==5.8.0
uvicorn==0.38.0
watchdog==6.0.0
wheel==0.44.0
xgboost==3.1.2
zstandard==0.19.0
```

Random Seed Initialization

- To ensure that all experiments in the project produce consistent and repeatable results, random seeds are initialized at the beginning of the workflow.
- Machine learning pipelines rely on randomness in several stages:
 - Train/validation splitting
 - Model weight initialization
 - Feature shuffling
 - Sampling
 - Any NumPy or Python-based random operations
- If the random seed is not fixed, you will get different results every time, which makes debugging and comparing models impossible.

Code Used for Random Seed Initialization

```
import random
import numpy as np

SEED = 42
random.seed(SEED)
np.random.seed(SEED)
```